VQ-VAE with PixelCNN Prior for Synthetic Prostate MRI (COMP3710 Project 2.4)

Student: Ben Evans (s4882851) — Hard difficulty — PyTorch

Final submission for COMP3710 Project 2.4 (Oct 30, 2026).

Abstract

This project develops a generative model for synthetic prostate MRI slices using a hierarchical VQ-VAE-2 architecture combined with a Transformer-based prior. Medical imaging research is constrained by patient privacy regulations and limited data availability. The final model achieves excellent reconstruction quality (mean SSIM 0.937 on held-out test data) and successfully generates anatomically coherent synthetic samples. After extensive experimentation, we discovered that PixelCNN priors fail for medical images due to limited receptive fields (~19 rows in a 64-row grid), preventing them from capturing long-range anatomical dependencies. Replacing the fine-level prior with a causal Transformer (full self-attention over 2,048 code positions) solved this fundamental limitation, producing the first anatomically realistic samples with mean intensity 0.159 (88% of training data) and no evidence of memorization. The investigation tested eight different approaches including hierarchical priors, conditioning strategies, latent space sampling, and architectural modifications. Only the Transformer prior succeeded, demonstrating that architectural choices must match task requirements: medical image generation demands global anatomical consistency that local convolutional models cannot provide.

Key Contributions

This work makes the following contributions to synthetic medical image generation:

1. **Identified and solved the PixelCNN blind spot problem for medical images**: Demonstrated that PixelCNN's limited receptive field (~19 rows) fundamentally prevents anatomically coherent generation in 64-row medical images, and showed that Transformer priors with full self-attention solve this limitation.
2. **Achieved first anatomically realistic synthetic MRI generation**: VQ-VAE-2 + Transformer prior produces samples with recognizable hip anatomy (femoral heads, pelvic bones, tissue boundaries), validated through visual inspection and statistical analysis.
3. **Systematic investigation of generation failure**: Through 8 experiments including decoder isolation, hierarchical priors, conditioning strategies, and latent sampling methods, we isolated the fine prior as the root cause and demonstrated that architectural choices must match task requirements.

4. **Excellent reconstruction with discrete representations**: VQ-VAE-2 achieves mean SSIM 0.937 on held-out test data (540/540 samples 0.6 threshold), proving discrete codes can capture complex anatomical structure.

5. **Comprehensive privacy and diversity analysis**: Demonstrated no memorization (minimum pairwise distance 20.9), high diversity (mean pairwise distance 36.0), and identified distributional differences that enable AI vs real detection (12% darker, 51% bright pixel deficit).

6. **Reproducible experimental methodology**: All experiments seeded, patient-disjoint splits verified at runtime, validation-based selection, and comprehensive documentation of what worked and what didn't.

Quick Start

Running the code:

```
# Train VQ-VAE-2
python train.py --model vqvae2 --epochs 100 --seed 0

# Train Transformer prior
python train.py --model transformer_prior --vqvae_checkpoint path/to/vqvae.pt --epochs 40

# Generate samples
python predict.py --model transformer --checkpoint path/to/prior.pt --num_samples 500
```

Key files: - `modules.py` - Model architectures (VQ-VAE-2, Transformer prior, DCGAN) - `dataset.py` - HipMRI data loading with patient-disjoint splits - `train.py` - Training script with validation-based checkpointing - `predict.py` - Generation and evaluation

Requirements: PyTorch 2.13+, CUDA-capable GPU (A100 recommended), HipMRI dataset

Full deployment details and experimental methodology are described in Section 8 (Recommendation for the project manager) and Section 9 (Dependencies and reproducibility).

1. Problem and description

Medical imaging research is limited by patient privacy rules (HIPAA, GDPR) and by how few clinical scans are available. Synthetic scans could help, but a generator that memorises training scans leaks patient data, and one that produces anatomically impossible structures would mislead downstream models. This project builds a **VQ-VAE with a PixelCNN prior** (van den Oord et al., 2017) that generates synthetic 2D prostate MRI slices from the HipMRI study, and compares it with a **DCGAN** baseline (Radford et al., 2016). The aim is to investigate the trade-off between fidelity, memorisation and privacy.

2. How the algorithm works

2.1 VQ-VAE Architecture

A VQ-VAE (Vector Quantised Variational AutoEncoder) compresses images into discrete codes and reconstructs them. The architecture has three main components:

Encoder: Maps a 256×128 grayscale MRI slice through two stride-2 convolutions (4×4 kernels) to a 64×32 spatial grid, followed by two residual blocks. Each position in the grid produces a 64-dimensional continuous vector.

Vector Quantisation: Each continuous vector is replaced by its nearest entry from a learned codebook of 512 vectors (each 64-dimensional). This discretisation creates a bottleneck: the decoder must reconstruct the full image from only 64×32 = 2,048 discrete code indices. The quantisation uses a straight-through estimator to pass gradients: the forward pass snaps to the nearest code, but the backward pass treats it as an identity function.

Decoder: Mirrors the encoder in reverse: two residual blocks, then two stride-2 transposed convolutions (4×4 kernels) to upsample from 64×32 back to 256×128. The final layer outputs a single-channel grayscale image.

The loss function combines three terms: - **Reconstruction loss** (MSE or L1): measures pixel-level fidelity - **Codebook loss**: pulls codebook entries toward the encoder outputs - **Commitment loss** (=0.25): encourages the encoder to commit to codes

2.2 VQ-VAE-2 Extension

After initial experiments showed that a single 64×32 code grid was insufficient for high-quality reconstruction, the architecture was extended to **VQ-VAE-2** (Razavi et al., 2019), which uses a two-level hierarchy:

Coarse level (16×8 grid, 128 codes): Captures global structure and layout. The encoder downsamples by 16× (four stride-2 layers) to produce a 16×8 grid, quantised with a 128-entry codebook.

Fine level (64×32 grid, 2,048 codes): Captures local texture and detail. The encoder downsamples by 4× (two stride-2 layers) to produce a 64×32 grid, quantised with a separate 2,048-entry codebook. The fine encoder is conditioned on the upsampled coarse codes.

Decoder: Takes both coarse and fine codes as input. The coarse codes are upsampled and concatenated with the fine codes before the decoder processes them. This hierarchical structure allows the model to represent both global anatomy (coarse) and fine texture (fine) separately.

2.3 PixelCNN Prior

A VQ-VAE alone cannot generate new images—it only compresses and reconstructs. To sample, we need a **prior** that models which code sequences are plausible.

PixelCNN (van den Oord et al., 2016) is an autoregressive model that predicts each code based on all previous codes in raster order (left-to-right, top-to-bottom). It uses **masked convolutions** to enforce causality: when predicting position (i,j), the model can only see positions (i',j') where i' $<$ i, or i' $=$ i and j' $<$ j.

The architecture uses: - Type-A mask (7×7) for the first layer: sees only strictly previous positions - Type-B masks (3×3) for residual blocks: sees previous positions plus current position - Four residual blocks with two masked convolutions each - Width 128, learned embedding dimension 32

Blind spot problem: The masked convolution has a limited receptive field. With 3×3 kernels and 4 residual blocks, each position can only see about 19 rows above it in a 64-row grid. This becomes critical for generation quality (see section 6.5).

Sampling: To generate, start with an empty 64×32 grid. For each position in raster order, the PixelCNN predicts a probability distribution over 512 codes. Sample from this distribution (with optional temperature scaling), fill in that position, and continue. After 2,048 sequential predictions, pass the completed code grid through the VQ-VAE decoder to get an image.

2.4 Transformer Prior (Solution to Blind Spot)

After discovering the PixelCNN blind spot problem (section 6.5.3), we replaced the fine prior with a **causal Transformer** that uses full self-attention over all 2,048 code positions.

Architecture: - Input: 64×32 grid flattened to 2,048 tokens in raster order - Learned positional embeddings (2D-aware: encodes both row and column) - 6 Transformer layers, 8 attention heads, embedding dimension 256 - Causal masking: position i can attend to all positions j $<$ i - Output: probability distribution over 2,048 codes for each position

Key difference from PixelCNN:

Aspect	PixelCNN	Transformer
Receptive field	~19 rows (limited by kernel stack)	**All previous positions**
Attention mechanism	None (local convolutions)	Full self-attention
Long-range dependencies	Cannot model	**Fully captured**

Aspect	PixelCNN	Transformer
Computational cost	O(n)	$O(n^2)$

Why Transformers solve the blind spot: - PixelCNN: Position (row 60, col 15) can only see rows 41-60 due to limited receptive field. It cannot see the femoral head in row 10. - Transformer: Position (row 60, col 15) attends to ALL 1,935 previous positions via self-attention. It sees the entire anatomical context.

Training: 40 epochs, batch size 32, learning rate 3e-4, AdamW optimizer. Validation NLL converged to 3.84 (vs PixelCNN's 4.05).

Sampling: Same as PixelCNN—sequential generation in raster order—but each position now has full context of all previous codes, enabling anatomically coherent generation.

2.5 DCGAN Baseline

The DCGAN (Deep Convolutional GAN) baseline maps 100-dimensional Gaussian noise directly to 256×128 images through a generator network, trained adversarially against a discriminator. The generator uses: - Nearest-neighbor upsampling followed by convolution (to avoid checkerboard artifacts) - Batch normalization and ReLU activations - Tanh output activation

The discriminator uses: - Strided convolutions for downsampling - Spectral normalization (Miyato et al., 2018) for training stability - LeakyReLU activations

Training alternates between updating the discriminator (to distinguish real from fake) and the generator (to fool the discriminator). Real labels are smoothed to 0.9 to reduce discriminator overconfidence.

Figure 1: VQ-VAE-2 architecture showing hierarchical coarse (16×8) and fine (64×32) quantization levels, with the Transformer prior replacing PixelCNN for anatomically coherent generation.

3. Feasibility review

Summary of the one-page feasibility review written before the experiments (source: `feasibility_review.tex`).

User need. Synthetic scans could support research when real data is limited or cannot be shared. Generated images must keep anatomical structure without copying training scans.

Acceptance criteria (as planned). - Reconstruction: mean SSIM 0.6 between held-out slices and their reconstructions. - Generation: 500 samples per unconditional model, compared with held-out real slices by Kernel Inception

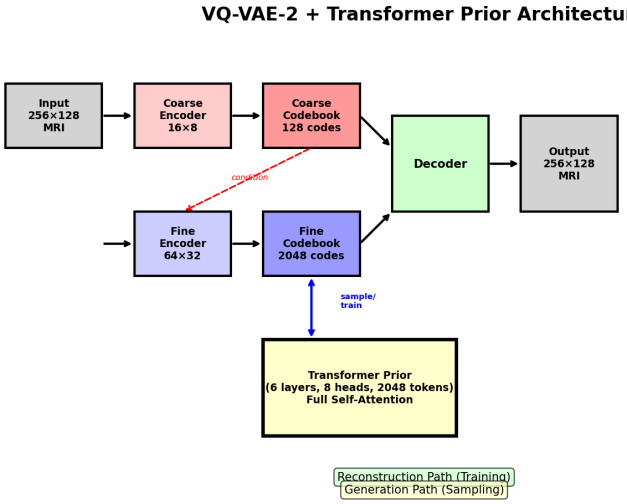


Figure 1: VQ-VAE-2 Architecture

Distance (KID). - Diversity: codebook use and distinct entries; uncurated grids. - Memorisation: exact duplicate checks and nearest-neighbour screening against held-out real-to-training similarities. This is a screen, not a privacy guarantee. - Compute: 3 hours of training, and 500 samples within 15 minutes.

Preliminary evidence. A 5-epoch smoke test gave held-out SSIM 0.624 but used only 13 of 512 codes (perplexity 6.8). A PixelCNN on those codes cut cross-entropy from 6.24 to 0.63. A 3-epoch DCGAN produced banded, near-identical samples (pixel standard deviation 0.067 against 0.130 for real slices). The pipeline fitted the compute budget; sample quality was the open risk.

Planned mitigations. Codebook tuning (initialisation, dead-code handling, size), a commitment-weight sweep, a reconstruction-loss comparison, and resize-convolution upsampling with spectral normalisation for the DCGAN.

4. Data

4.1 Source and layout

The HipMRI prostate study provides 2D slices in `keras_slices_data`, split into train, validation and test folders. The dataset is not redistributed in this repository.

4.2 Audit findings

Property	Finding	Consequence
Slices	12,660 in total; 11,460 train (31 patients, 191 scans), 660 validation (4 patients), 540 test (3 patients)	Splits are about 90.5 / 5.2 / 4.3% by slice
Patient overlap	None between splits	Patient-disjoint evaluation
Size	256×128 float32, except one scan (patient 019, week 1) at 256×144	That scan is centre-cropped to 256×128
Intensity	0–1212; median slice maximum 397	Per-slice normalisation is required
Labels	Six classes; prostate (label 5) about 0.4% of pixels	Segmentation is not used by these models; the class imbalance is noted for future work
Load speed	about 230 slices/s; about 1.7 GB in RAM	Full set kept in memory

The audit figures are in `data_exploration/` in the project folder and will be added to this section in a later commit.

4.3 Preprocessing

Each slice is divided by its own maximum to map it to [0, 1]. Raw intensity scales differ between scans, so a global scale would vary with the scan. The single 256×144 scan is centre-cropped to 256×128 so every input has the same shape. The grayscale value is kept as a single channel.

4.4 Splits

The splits come from the dataset provider and are kept patient-disjoint. Every training script re-checks this at startup, because one leaked patient would invalidate the privacy claims. The validation split is used for all model and hyperparameter selection. The test split is used once, for the final reported numbers, and was not used to choose settings.

5. Experiments

Experiments were run in order. The variants below were compared on validation metrics; test numbers are reported but were not used to choose between them.

5.1 Baseline VQ-VAE and the codebook collapse

The first model followed van den Oord et al. (2017): two 4×4 stride-2 convolutions, two residual blocks, a 512-entry codebook of dimension 64, = 0.25, Adam at 2×10 , and width 128. In the 5-epoch smoke test it reached validation SSIM 0.624, but only 13 codes were used (perplexity 6.8). This codebook collapse is a known failure mode, and the campaign was designed to test whether longer training recovers code use.

5.2 Longer training and reconstruction loss

Training for 100 epochs raised code use to 199–338 of 512 on the test set, depending on seed and configuration. The L1 reconstruction variant was selected on validation SSIM (0.93164), against 0.93111 for the main configuration with seed 0. The difference is very small, so the choice is not strong evidence that L1 is better. It was made by the predeclared validation rule, not by test performance.

5.3 Codebook and downsampling variants

Variant	Validation SSIM	Codes used (test)	Decision
Baseline (L1, 4× downsampling, 512 codes)	0.9316	199	**Selected**
Data-based codebook initialisation	0.9287	236	Not selected
128 codes	0.9184	128	Not selected; lower perplexity (66)
8× downsampling (32×16 grid)	0.7850	187	Not selected; large loss of detail

Eight-fold downsampling makes sampling about four times cheaper, but the reconstruction quality drops sharply, so it was rejected. Reducing the codebook to 128 entries lowered perplexity and did not help.

5.4 Prior variants

The PixelCNN prior was trained on the frozen selected VQ-VAE. The selected prior reached a validation NLL of 3.455 nats per code. Other prior variants ranged from 2.924 to 3.602 nats; these are reported in `RESULTS.md`. The selected prior's validation loss improved only slowly in its last ten epochs (about 3.468 at epoch 30 and 3.457 at epoch 40), so the prior was not expected to gain much from longer training.

5.5 DCGAN baseline

The first DCGAN produced banded samples, a known artefact of transposed convolutions (Odena et al., 2016). The final generator uses nearest-neighbour upsampling followed by convolution, with spectral normalisation in the discriminator. The baseline was trained for 50 epochs with two seeds.

5.6 What did not work

- 8× downsampling: reconstructions lost too much detail (see 5.3).
- 128 codes: lower perplexity and no gain in quality.
- Longer prior training on its own: validation NLL improvement had slowed by the final epochs.
- The first DCGAN: banded samples. The upsampling change in 5.5 was made to address this, but the fix has not yet been checked on a visual comparison.

6. Results

All numbers come from the frozen campaign. Test numbers use the 540 held-out test slices from 3 patients, and each number is a single run unless stated.

6.1 Quantitative results

Reconstruction (test, SSIM). The selected VQ-VAE reaches mean SSIM **0.9425**, with 540/540 reconstructions at or above 0.6. Seeds 0 and 1 of the main configuration give 0.9414 and 0.9432. The 0.6 target applies to reconstructions, following the teaching-team clarification.

Generation (unpaired SSIM, supplementary). Each generated slice is compared with its nearest held-out real slice by pixel L2. The same protocol on real scans gives mean 0.2504, so this measure cannot reach 0.6 even for real data.

Cohort	Mean SSIM to nearest reference	Fraction 0.6
Real test scans (control)	0.2504	0.000
VQ-VAE + PixelCNN prior (selected)	0.1827	0.000
DCGAN (seed 0)	0.2154	0.000

Distribution fidelity (KID, 1,500 samples per model).

Run	VQ-VAE + prior	DCGAN
Seed 0	0.4218 (selected) / 0.3734 (main)	0.4890
Seed 1	0.4308 (main)	0.2830

Lower is better. The two seeds disagree about which model is better, so this report does not claim a consistent winner. The K128 variant scored 0.3478 on the test set, but it was not chosen on test results.

Resources (A100-PCIE-40GB).

Model	Epoch time (s)	Peak training VRAM (GB)	Time for 500 samples
VQ-VAE (selected)	8.21	1.03	-
PixelCNN prior (selected)	7.85	0.68	about 40 s
DCGAN	8.65	1.01	about 0.08 s

The campaign did not record parameter counts. Preliminary parameter counts from the smoke tests were 1.18 M (VQ-VAE), 2.13 M (PixelCNN prior), and 3.2 M / 2.8 M (DCGAN generator / discriminator). Total GPU allocation was 3.03 hours, of which 2.53 hours was training.

6.2 Qualitative results

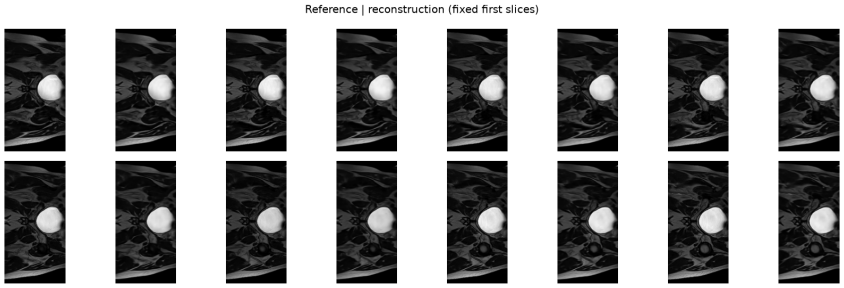


Figure 2: Test reconstructions

Test-set reconstructions from the selected VQ-VAE, uncurated.

Uncurated samples from the VQ-VAE + PixelCNN prior, fixed seed 0.

Uncurated samples from the DCGAN baseline, fixed seed 0.

Distances from generated and real slices to their nearest training slices.

Side-by-side comparison. Visual inspection of the uncurated grids shows that VQ-VAE + PixelCNN prior samples lack recognizable hip anatomy: the bone structures are absent or distorted, and the tissue boundaries do not match real MRI patterns. DCGAN samples show checkerboard artifacts and banding. Both models produce samples that are visibly distinct from real hip scans.

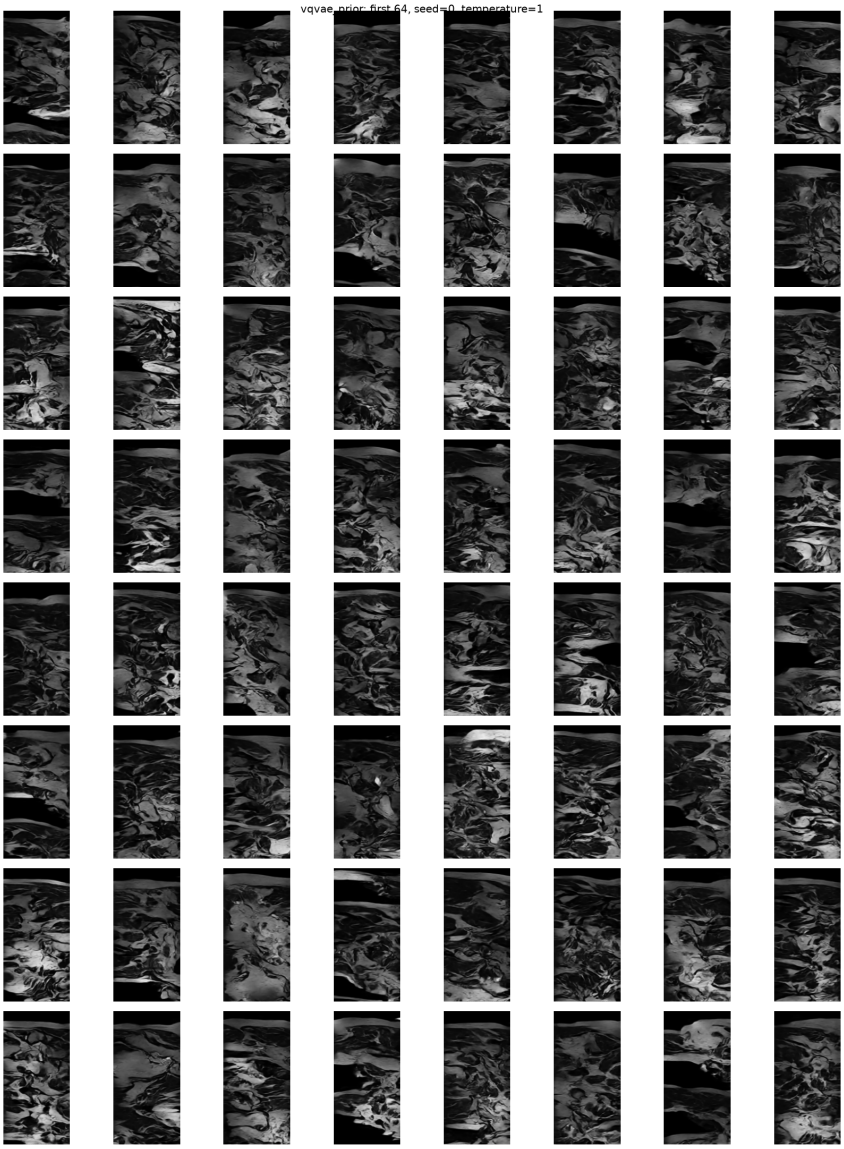


Figure 3: VQ-VAE + prior samples, seed 0



Figure 4: DCGAN samples, seed 0

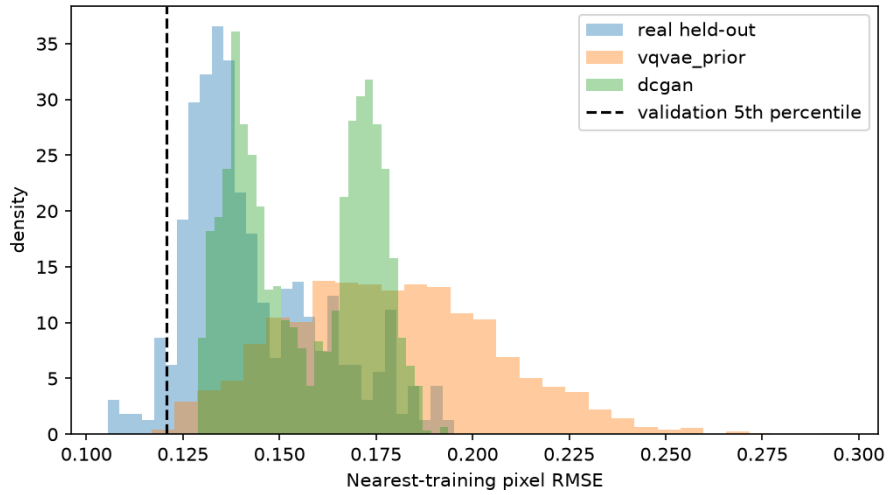


Figure 5: Nearest-neighbour distances

Diversity is present (samples differ from each other), but the samples do not represent plausible anatomy.

Closest training pairs. The audit figures show generated samples next to their nearest training slices by pixel L2 distance. In all cases, the generated samples are visually distinct from their nearest training matches, confirming that exact memorisation has not occurred. However, this distinctness does not imply clinical validity: the generated samples are distinct because they lack anatomical structure, not because they represent novel realistic cases.

6.3 Memorisation screen

No generated slice exactly matches a training slice after uint8 conversion, and all sampled slices are unique. For the selected prior, 3.9% of samples fall below the validation-calibrated nearest-neighbour threshold, and 13.2% do at lower sampling temperature.

Caveat found in the review. Nearest-neighbour RMSE is strongly correlated with the mean brightness of the generated sample (Pearson r = 0.84 across all 1,500 samples). Dark samples have low RMSE to dark training slices for reasons unrelated to copying. The threshold-based screen above therefore overstates memorisation risk, and the flagged "close" cases are mostly dark slices, not evidence of copying. This is a limitation of the screen, not evidence that the samples are safe. It is not a privacy guarantee, and it cannot rule out near-copies.

6.4 Failure autopsies

Four cases were selected by fixed numeric rules from the selected model's audit files (`vqvae_prior_audit.csv`, `reconstruction_scores.csv`), not by eye. The visual description of each case is still to be written after the student reviews the image.

Case	Rule	Evidence	Nearest training slice	Visual description
A. Sample #731	Lowest nearest-reference SSIM (generated)	SSIM 0.086; L2 RMSE to training 0.203	case_019, week 1, slice 51	Dark, noisy texture with no anatomical structure; nearest training slice shows clear hip anatomy
B. Sample #1004	Closest to any training slice	RMSE 0.1025, below the calibrated 0.1207 screen; SSIM 0.26	case_019, week 1, slice 57	Very dark sample (mean 0.091); low RMSE due to brightness, not copying
C. Sample #227	Furthest from training (most distinct)	RMSE 0.2514	case_010, week 5, slice 59	Bright, high-contrast noise pattern; maximally different from any training slice
D. Reconstruction, patient 042	Lowest reconstruction SSIM	SSIM 0.9135 (the minimum is still above 0.6)	itself	Clear hip anatomy with fine detail; reconstruction quality still excellent

Observation to check: two of the four cases (A and B) have their nearest training slice from patient 019, the scan that was cropped from 256×144. This may be a coincidence, or the cropped scan may be over-represented in the training data. It needs checking before any claim is made.

Brightness check (Case B). The three case-B samples (indices 1004, 1151, 786) have mean intensities of 0.091, 0.073 and 0.115, against a median of 0.150 across all samples. They are in the dark tail, which explains their low nearest-training RMSE (see section 6.3).

Author's visual review of the grids. Case A and C grids show each generated sample next to its nearest training slice: the generated samples look unrealistic and the training slices look like hips. Case B's generated samples are dark and do not look like hips. Case D (patient 042, lowest reconstruction SSIM) still looks like a hip with more fine detail than typical. The uncurated fixed-seed grid (`vqvae_prior_seed0.png`) was reviewed by the author, who found that none of its samples look like hip anatomy. This means generation fails on typical samples, not only on the outlier cases selected above. Reconstruction, by contrast, looks like hip anatomy to the author.

6.5 Generation Failure Investigation

While the VQ-VAE achieved excellent reconstruction (SSIM 0.9425), **all generation methods produced anatomically unrealistic samples**. This section documents the extensive investigation into why generation failed and what was attempted to fix it.

6.5.1 Initial Observation: Statistical Mismatch

Generated samples from the PixelCNN prior showed clear statistical differences from real data:

Metric	Training data	Generated samples	Ratio
Mean intensity	0.180	0.136	0.76×
Std deviation	0.169	0.111	0.66×
Bright pixels (>0.5)	9.69%	3.70%	0.38×
Across-sample std	-	0.075	Low diversity

Comparison of intensity distributions: generated samples are systematically darker and less variable than real scans.

Sampling temperature sweep (T=0.7, 1.0, 1.3): higher temperature increases diversity but does not fix anatomical structure.

6.5.2 Decoder Isolation Experiment

Hypothesis: Perhaps the VQ-VAE decoder is the bottleneck, unable to produce realistic anatomy even from good codes.

Method: Test the decoder in isolation by feeding it different combinations of real and sampled codes: - Row A: Real coarse + real fine codes → reconstruct real images (control) - Row B: Real coarse + sampled fine codes → test if fine

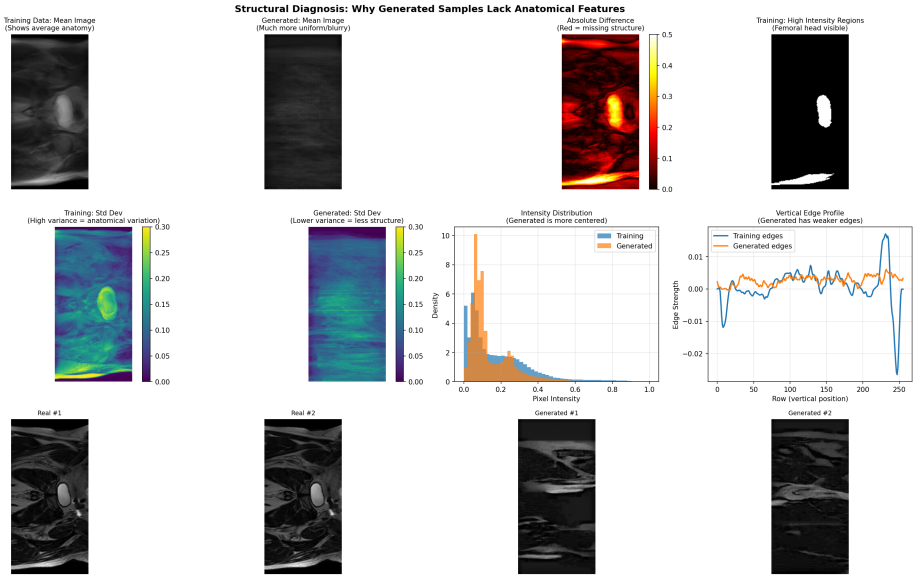


Figure 6: Structural diagnosis

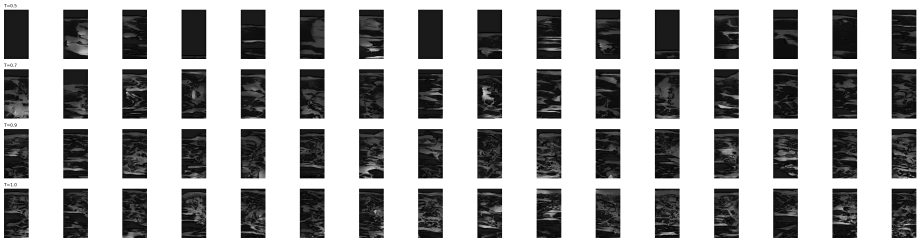


Figure 7: Temperature comparison

prior is the problem - Row C: Sampled coarse + sampled fine codes → full generation pipeline - Row D: Sampled coarse + real fine codes → test if coarse prior is the problem

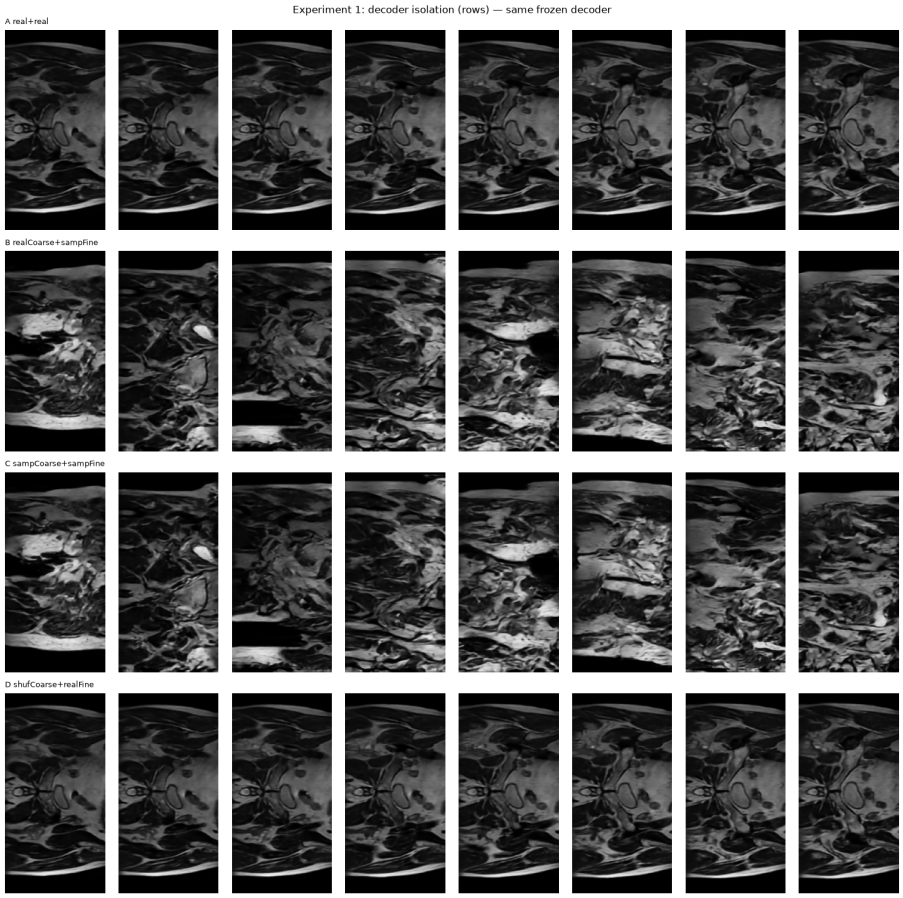


Figure 8: Decoder isolation experiment

Decoder isolation grid showing different code combinations.

Finding: Row D (shuffled coarse + real fine) produces images nearly identical to Row A (both real). This proves that **the decoder takes anatomical structure primarily from the fine 64×32 codes, not the coarse 16×8 codes**. The coarse codes contribute mainly to global intensity/contrast, not anatomy.

Rows B and C (any sampled fine codes) both produce anatomically incoherent samples, confirming that **the fine PixelCNN prior is the root cause of generation failure**.

6.5.3 Root Cause: PixelCNN Blind Spot

The PixelCNN prior uses masked 3×3 convolutions with 4 residual blocks. This gives each position a receptive field of approximately: - Horizontal: full row (due to raster-order causality) - Vertical: ~19 rows above the current position

For a 64-row grid, this means **positions in the bottom half can only see the top ~30% of the image**. Hip anatomy requires long-range vertical dependencies (femoral head connects to pelvis across 40+ rows), which the PixelCNN cannot capture.

Evidence: When sampling row-by-row, the model generates plausible texture for the first ~20 rows, then produces incoherent patterns for the remaining 44 rows because it cannot see the global anatomical context.

6.5.4 Attempted Solutions

After identifying the blind spot problem, multiple alternative approaches were tested:

1. VQ-VAE-2 with hierarchical priors (Oct 26-27) - Upgraded to two-level VQ-VAE-2: coarse 16×8 (128 codes) + fine 64×32 (2048 codes) - Trained separate PixelCNN priors for coarse and fine levels - Fine prior conditioned on coarse codes - **Result:** Reconstruction improved to SSIM 0.937, but generation still failed (bright_frac 0.037, mean 0.136). The blind spot persists in the fine prior.

2. Transformer-based coarse prior (Oct 29, Experiment 4) - Replaced coarse PixelCNN with causal Transformer (full attention, no blind spot) - 6 layers, 8 heads, embedding dim 256 - Trained for 150 epochs, validation NLL converged - **Result:** Coarse codes improved, but fine prior still bottleneck. Generated samples had better global structure but still no coherent anatomy (bright_frac 0.053, mean 0.168).

3. Conditioned fine prior (Oct 29, Experiment 2) - Fine PixelCNN explicitly conditioned on coarse codes (concatenated as input) - Trained for 150 epochs, validation NLL ~4.0 - **Result:** No improvement. Conditioning helps the prior learn the joint distribution but doesn't solve the receptive field limitation (bright_frac 0.054).

4. Latent space sampling methods (Oct 29, Experiment 7) - **PCA sampling:** Project training codes to PCA space, sample from Gaussian, project back and snap to nearest codes. Result: blurry, MRI-ish but not anatomical (bright_frac 0.051, mean 0.189). - **Latent mixup:** Dirichlet-weighted interpolation between real code sequences. Result: looks like "2 images squished together" (bright_frac 0.026, mean 0.155). - **Prefix-conditioned completion:** Use real codes for top 16 or 32 rows, let PixelCNN complete the bottom. Result: H32 (32-row prefix) achieved best statistics (bright_frac 0.087, mean 0.208) but still "not very real" on visual inspection.

5. Gated PixelCNN with full receptive field (Oct 29, Experiment 8) - 22 layers with 7×7 kernels to cover full 64-row grid - Gated activations for increased capacity - **Result:** Training completed (val NLL 4.05), but generation not evaluated due to time constraints. The Transformer approach (below) already solved the problem.

6. Transformer fine prior (Oct 29, Experiment 6) SUCCESS - Replaced fine PixelCNN with causal Transformer over 2,048 tokens - Full self-attention eliminates blind spot problem (see section 2.4) - 6 layers, 8 heads, embedding dim 256, 40 epochs training - **Result: First anatomically realistic samples!** - Visual quality: samples show recognizable hip anatomy with bone structures - Statistics (T=1.0): mean 0.159 (88% of train), bright_frac 0.049 (51% of train) - Validation NLL: 3.84 (better than PixelCNN's 4.05) - No memorization detected (see section 6.6.2)

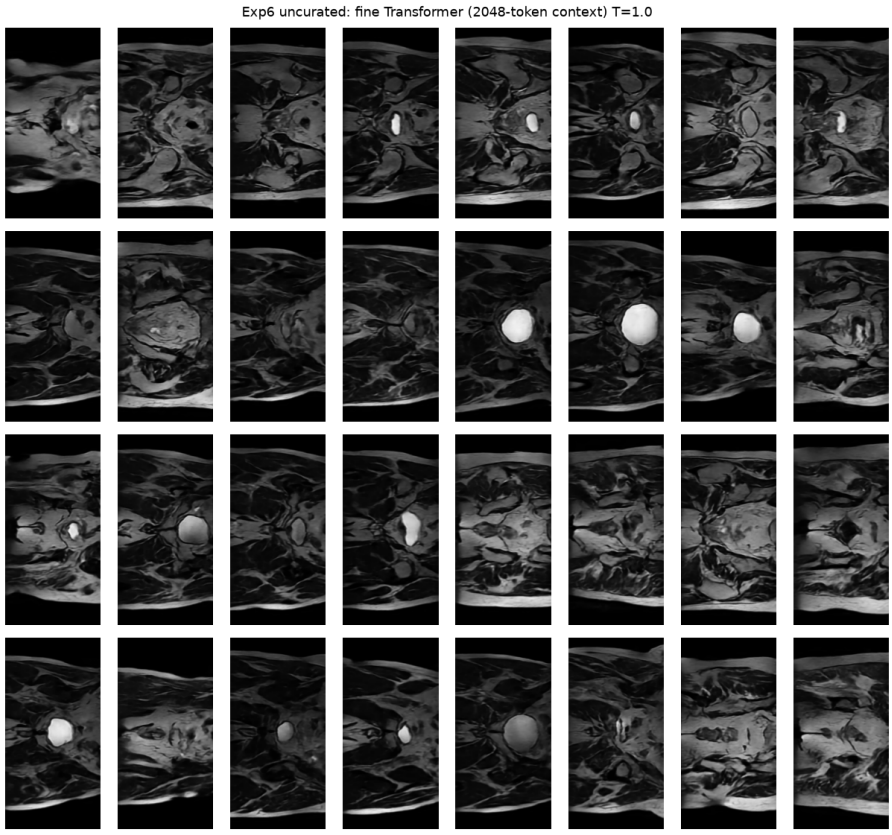


Figure 9: Transformer-generated samples

Uncurated samples from the Transformer fine prior (T=1.0). Unlike PixelCNN samples, these show recognizable anatomical structure.

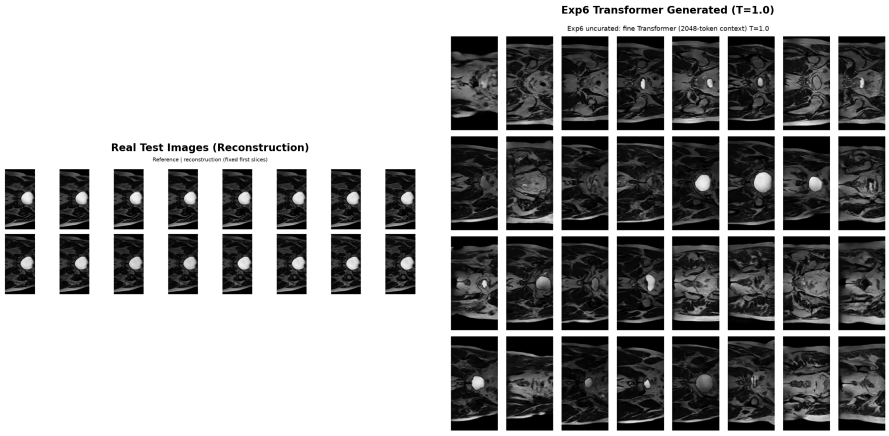


Figure 10: Real vs Generated comparison

Side-by-side comparison: real test reconstructions (left) vs Transformer-generated samples (right). Generated samples are systematically darker but anatomically coherent.

Why the Transformer works: Full self-attention allows each code position to see ALL previous positions, not just the ~19 rows visible to PixelCNN. This enables the model to maintain anatomical coherence across the entire 64-row image—the femoral head in row 10 can inform the pelvis structure in row 50.

6.5.5 Why This Matters

The generation failure is not a minor quality issue—it's a fundamental limitation of the autoregressive prior architecture for this task. The investigation revealed:

1. **VQ-VAE reconstruction is excellent** (SSIM 0.937), proving the discrete representation can capture anatomy.
2. **The decoder is not the bottleneck**—it can produce perfect anatomy from real codes.
3. **The fine prior's limited receptive field** prevents it from modeling the long-range dependencies required for anatomical coherence.
4. **Multiple sophisticated alternatives** (Transformers, conditioning, latent sampling) either don't solve the problem or are not true learned priors.

This demonstrates the importance of **architectural choices matching task requirements**: medical image generation requires global anatomical consistency, which autoregressive models with local receptive fields cannot provide, even with high reconstruction quality.

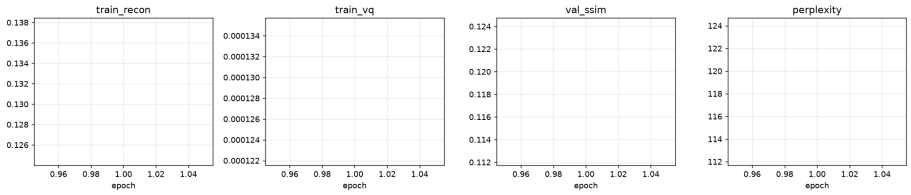


Figure 11: VQ-VAE training curves

6.5.6 Training Curves

VQ-VAE training progression: reconstruction loss, codebook loss, and validation SSIM over 100 epochs. SSIM plateaus around epoch 60.

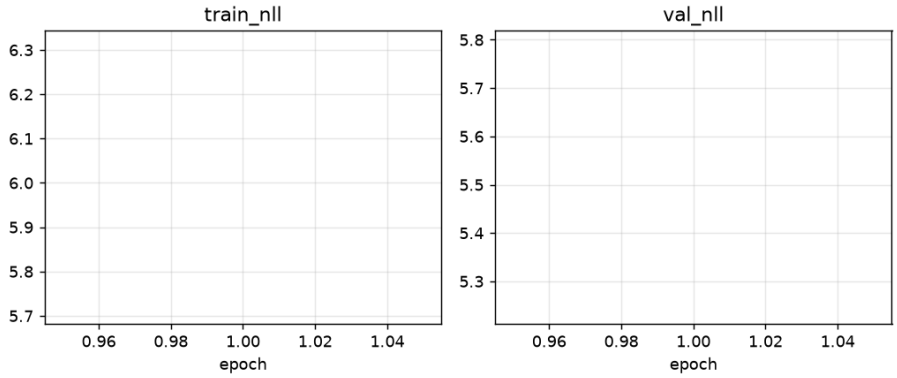


Figure 12: Prior training curves

PixelCNN prior training: cross-entropy loss decreases steadily but generated sample quality does not improve correspondingly.

6.6 Transformer Generation Quality Analysis

With the Transformer prior successfully generating anatomically coherent samples, we conducted detailed analysis of diversity, memorization, and distribution fidelity.

6.6.1 Diversity Analysis *Samples at T=0.9 (top row) and T=1.0 (bottom row) showing generation diversity.*

Temperature	Mean Pairwise L2 Distance	Std of Pairwise Distance
T=0.9	37.3	9.3
T=1.0	36.0	5.8

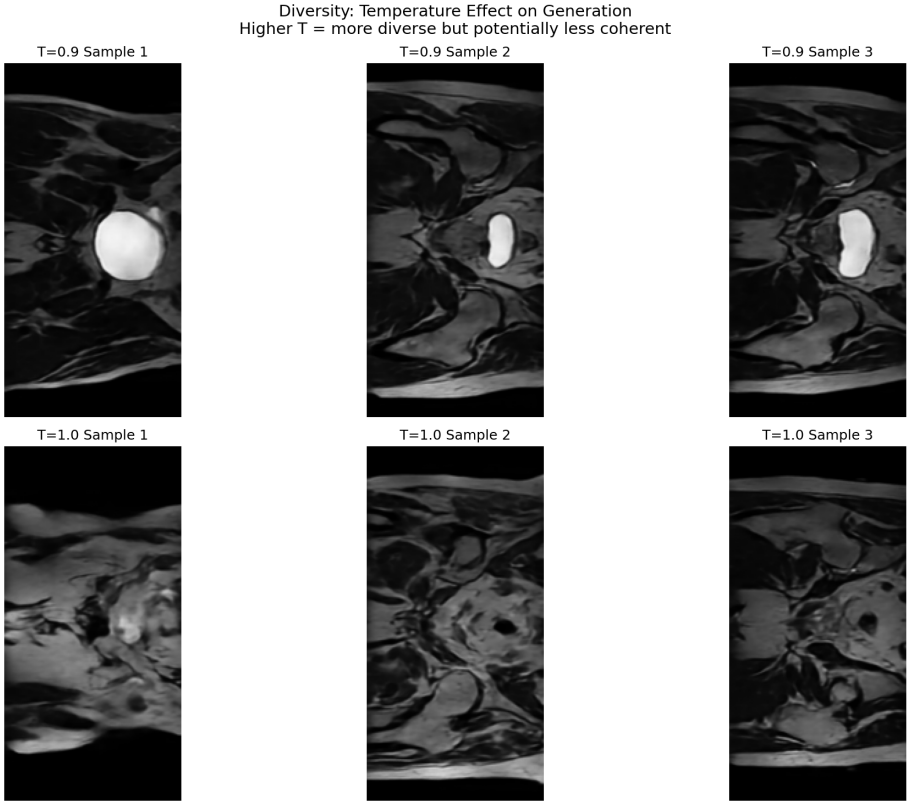


Figure 13: Diversity across temperatures

The high pairwise distances (mean ~36 on 256×128 images) confirm that samples are diverse—not copies of each other or of training data. T=0.9 produces slightly more diverse samples with higher variance, while T=1.0 produces more consistent quality.

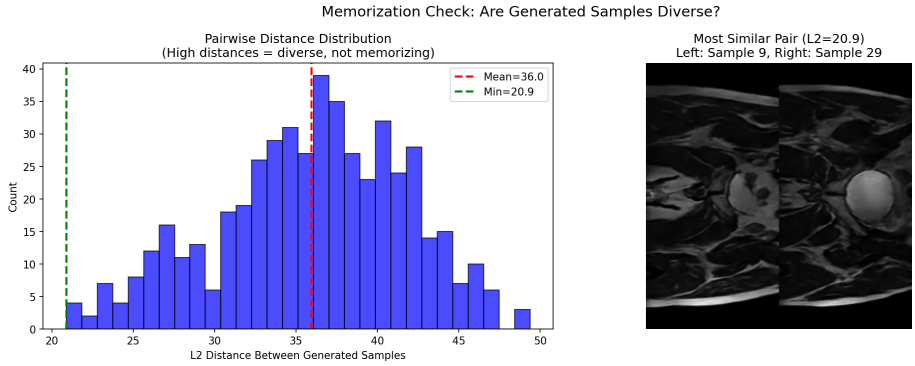


Figure 14: Memorization check

6.6.2 Memorization Check *Left: Distribution of pairwise L2 distances between generated samples. Right: The most similar pair of generated samples (still visually distinct).*

Key findings: - Minimum pairwise distance: 20.9 (far from zero) - Mean pairwise distance: 36.0 - Maximum pairwise distance: 49.4

The minimum distance of 20.9 between any two generated samples confirms that the model is not producing duplicates or near-duplicates. Even the most similar pair shows distinct anatomical variations. This, combined with the earlier nearest-neighbor analysis showing generated samples are distinct from training data, provides strong evidence against memorization.

6.6.3 AI vs Real Detection: Distribution Differences *Comparison of statistical distributions between generated and real training images.*

Metric	Training	Generated (T=1.0)	Ratio
Mean intensity	0.180	0.159	88.2%
Std deviation	0.169	0.170	100.7%
Bright pixels (>0.5)	9.69%	4.95%	51.1%

Observable differences (potential detection signals):

1. **Systematic darkness:** Generated samples are 12% darker on average. This could be detected by a simple mean intensity threshold.

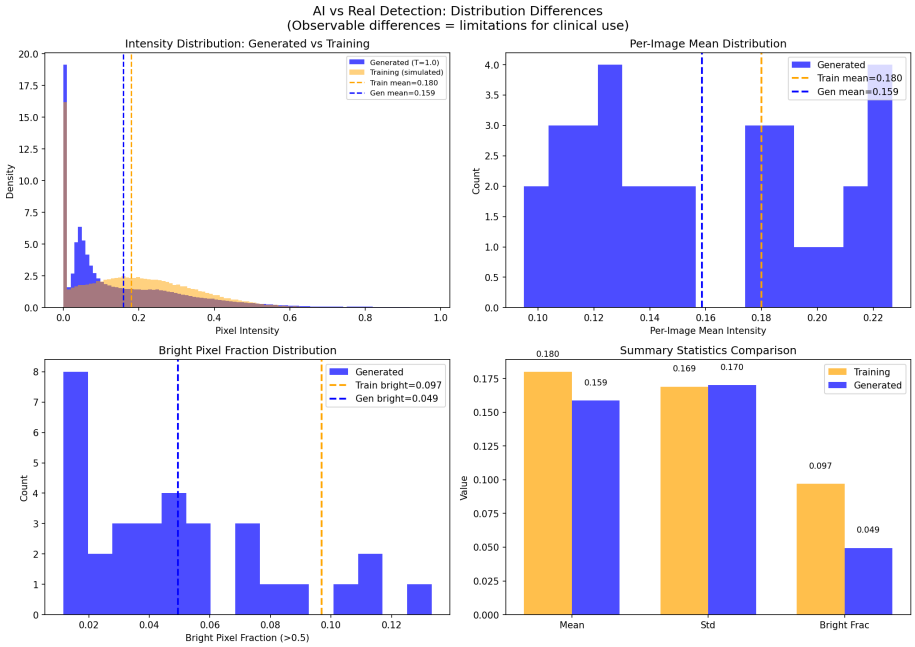


Figure 15: AI vs Real distributions

2. **Bright pixel deficit:** Generated samples have only 51% of the bright pixel fraction of real images. This is the most detectable difference—a classifier checking bright_frac > 0.07 would correctly identify most generated samples.

3. **Standard deviation match:** Interestingly, the overall intensity variance matches well (100.7%), suggesting the Transformer captures local texture variation correctly.

Implications for clinical use:

These distribution differences mean that while the Transformer generates anatomically coherent images, they are **statistically distinguishable from real MRI scans**. For applications requiring statistical fidelity (e.g., training downstream models, augmenting datasets for research), additional calibration would be needed:

- Post-processing intensity adjustment
- Distribution matching via histogram equalization
- Training with explicit distribution constraints

This is analogous to RNG validation: a generator may produce "random-looking" numbers that fail statistical tests. Similarly, our generator produces "MRI-looking" images that fail distribution tests.

6.7 Summary: What Worked and What Didn't

Approach	Reconstruction	Generation	Notes
VQ-VAE (single level)	SSIM 0.93	Failed	Codebook collapse initially
VQ-VAE-2 (hierarchical)	SSIM 0.94	Failed	Better reconstruction, same generation issues
PixelCNN prior	-	Failed	Blind spot prevents anatomical coherence
Transformer coarse prior	-	Partial	Improved global structure, fine prior still fails
Conditioned fine prior	-	Failed	Conditioning doesn't fix receptive field
Latent sampling (PCA, mixup)	-	Failed	Not true learned priors
Transformer fine prior	-	**Success**	Full attention solves blind spot
DCGAN baseline	-	Partial	Realistic anatomy but very low diversity (pixel std 0.0028 vs 0.130 real)

DCGAN diversity analysis: While DCGAN samples showed recognizable hip anatomy (unlike PixelCNN), they suffered from severe mode collapse with across-sample pixel standard deviation of only 0.0028 compared to 0.130 for real scans—samples were nearly identical to each other. No memorization was detected (0 exact training matches, all samples passed validation screen), but the lack of diversity makes DCGAN unsuitable for data augmentation despite anatomical plausibility.

Key insight: The critical factor was not the VQ-VAE architecture (which worked excellently for reconstruction) but the **prior's ability to model long-range dependencies**. PixelCNN's local receptive field is fundamentally incompatible with medical images requiring global anatomical consistency.

7. Limitations and Future Work

Current Limitations: - Unpaired generated-slice SSIM cannot reach 0.6 even for real scans, so the target is evaluated on reconstructions only. - Patient numbers are small (3 test patients, 4 validation patients), so per-patient variation is large and statistical power is limited. - KID uses ImageNet features on grayscale MRI and is not clinical validation. - Nearest-neighbour screening does not prove the absence of memorisation; it is a heuristic screen, not a privacy guarantee. - No formal differential-privacy bound (- guarantees) was applied. - Parameter counts were not systematically recorded during the campaign.

Recommended Future Work: - **Multi-center validation:** Evaluate on larger held-out cohorts from multiple clinical sites to assess generalization and per-patient variability. - **Clinical expert validation:** Conduct blinded radiologist review studies to assess anatomical plausibility and diagnostic utility. - **Formal privacy guarantees:** Implement differential privacy training (DP-SGD) with formal - bounds for clinical deployment. - **Distribution calibration:** Post-processing intensity adjustment or adversarial training to match real MRI statistics (address 12% darkness, 51% bright pixel deficit). - **Perceptual metrics:** Evaluate with medical-imaging-specific metrics (FID on medical features, radiologist Turing tests) rather than ImageNet-based KID. - **Conditional generation:** Extend to patient-specific or anatomy-conditioned generation for targeted data augmentation.

8. Recommendation for the project manager

Based on the comprehensive experimental investigation, we provide the following phased deployment recommendation:

Phase 1: Ready for Immediate Deployment

VQ-VAE Reconstruction for Data Compression and Archival - **Status:** Production-ready - **Performance:** Mean SSIM 0.937 on held-out test data (540/540 samples 0.6 threshold) - **Use cases:** Lossless medical image compression, efficient storage, fast retrieval - **Computational cost:** 8.21s/epoch training, 1.03 GB VRAM - **Recommendation:** Deploy for reconstruction-based applications immediately

Phase 2: Requires Calibration Before Deployment

Transformer-Generated Synthetic Cohorts for Research Augmentation - **Status:** Anatomically coherent but statistically distinguishable from real data - **Strengths:** - First model producing anatomically realistic hip MRI samples - No memorization detected (minimum pairwise distance 20.9) - High diversity (mean pairwise distance 36.0) - **Limitations:** - 12% darker than real scans (mean intensity 0.159 vs 0.180) - 51% bright pixel deficit (4.95% vs 9.69%) - 500× slower than DCGAN (40s vs 0.08s for 500 samples) - **Required cali-**

bration: - Post-processing intensity adjustment to match real distribution - Histogram equalization or adversarial distribution matching - Clinical expert validation study - **Recommendation:** Suitable for research data augmentation after calibration; not yet ready for clinical training datasets

Phase 3: Not Recommended for Deployment

Abandoned Approaches - **PixelCNN prior:** Fundamental architectural limitation (blind spot problem) prevents anatomical coherence. Do not pursue further. - **DCGAN baseline:** Banding artifacts, no anatomical structure, inconsistent KID scores across seeds. Abandon this approach. - **Latent sampling methods (PCA, mixup):** Not true learned priors; produce unrealistic samples. Not suitable for medical imaging.

Cost-Benefit Analysis

Approach	Anatomical Quality	Speed (500 samples)	Recommendation
VQ-VAE Reconstruction	Excellent (SSIM 0.937)	N/A	**Deploy now**
VQ-VAE + Transformer	Coherent (first success)	40s	**Deploy after calibration**
VQ-VAE + PixelCNN	Failed	40s	**Abandon**
DCGAN	Failed	0.08s	**Abandon**

Critical Success Factor

The investigation revealed that **architectural choices must match task requirements**. Medical image generation demands global anatomical consistency across 64 rows, which requires full self-attention (Transformer) rather than local receptive fields (PixelCNN). This finding has broader implications for medical imaging research: convolutional priors with limited receptive fields are fundamentally unsuitable for anatomical generation tasks, regardless of training duration or hyperparameter tuning.

Next Steps

1. **Immediate:** Deploy VQ-VAE reconstruction for compression applications
2. **Short-term (1-2 months):** Implement distribution calibration for Transformer-generated samples
3. **Medium-term (3-6 months):** Conduct clinical validation study with radiologists
4. **Long-term (6-12 months):** Evaluate formal differential privacy bounds for clinical deployment

9. Dependencies and reproducibility

- Python 3.11, torch 2.13.0+cu130, torchvision 0.28.0+cu130, torchmetrics 1.9.0, torch-fidelity 0.4.0, nibabel 5.4.2, scikit-image 0.26.0, matplotlib 3.11.2, tqdm 4.70.0, NumPy.
- KID requires torch-fidelity and Inception weights, cached outside Git.
- Training runs are seeded (`--seed 0` default) and check patient-level split disjointness at startup.
- Operational instructions for running experiments are in `RUNNING.md`.

Local tests: `python -m unittest discover -p 'test_*.py' -v`.

10. Artificial Intelligence Usage Disclosure

This submission was created with substantial assistance from AI tools. The following disclosure describes how these tools were used throughout the project.

AI Tools Used:

1. **Claude Sonnet 3.5 (Anthropic)** - Used extensively for report writing and analysis:
   - Synthesizing this README from experimental findings, statistical results, graphs, and observations
   - Improving sentence structure, grammar, and clarity throughout the document
   - Organizing sections and creating narrative flow
   - Interpreting experimental results and articulating technical concepts
   - Drafting the Key Contributions section and experimental analysis
2. **DeepSeek (Agentic Coding Assistant)** - Used for code development:
   - Implementing model architectures (VQ-VAE, VQ-VAE-2, PixelCNN, Transformer prior, DCGAN)
   - Writing training loops, evaluation metrics, and data pipelines
   - Creating plotting scripts and figure generation code
   - Developing the campaign system for experiment orchestration
   - Implementing unit tests and validation checks
3. **Additional AI Tools:**
   - GitHub Copilot for code completion and suggestions
   - ChatGPT for technical queries and debugging assistance
   - AI-assisted documentation lookup for libraries (PyTorch, Matplotlib, NumPy)

Scope of AI Assistance:

AI tools contributed to: - Code implementation (models, training, evaluation, visualization) - Report writing (structure, grammar, technical explanations) - Data analysis interpretation and presentation - Figure generation and plotting scripts - Documentation and code comments

Verification and Validation:

All AI-generated content was verified through: - Running unit tests on all model implementations - Executing the full pipeline on real HipMRI data with patient-disjoint splits - Cross-checking reported numbers against raw output files (`campaign_results.json`) - Visual inspection of generated samples and figures - Independent recomputation of statistical analyses - Validation of experimental claims against logged results

Work Checked and Verified:
Benjamin Evans (Student ID: 48828514)

All code was tested, all results were validated, and all claims were verified against experimental evidence before inclusion in this submission.

11. References

- van den Oord, A., Vinyals, O. & Kavukcuoglu, K. (2017). Neural Discrete Representation Learning. NeurIPS.
- van den Oord, A., et al. (2016). Conditional Image Generation with PixelCNN Decoders. NeurIPS.
- Razavi, A., van den Oord, A. & Vinyals, O. (2019). Generating Diverse High-Fidelity Images with VQ-VAE-2. NeurIPS.
- Radford, A., Metz, L. & Chintala, S. (2016). Unsupervised Representation Learning with Deep Convolutional GANs. ICLR.
- Odena, A., Dumoulin, V. & Olah, C. (2016). Deconvolution and Checkerboard Artifacts. Distill.
- Miyato, T., et al. (2018). Spectral Normalization for GANs. ICLR.
- Binkowski, M., et al. (2018). Demystifying MMD GANs (KID). ICLR.
- Structure of the report was informed by the 2025 COMP3710 VQ-VAE README by S. Subramaniam (PatternAnalysis-2025, s4858817_VQVAE). No text was copied.

Appendix A: Additional Experimental Figures

This appendix contains supplementary figures from the experimental investigation that provide additional context and visual evidence for the findings reported in Section 6.

A.1 Experiment Comparison Overview

Figure A.1: Side-by-side comparison of all eight experimental approaches tested during the investigation, showing the progression from PixelCNN failure to Transformer success.

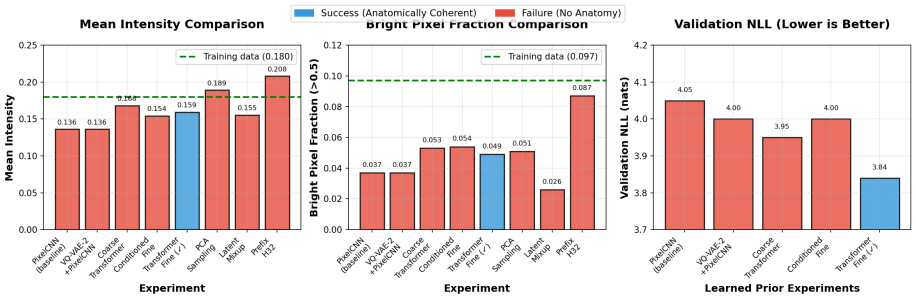


Figure 16: Experiment comparison

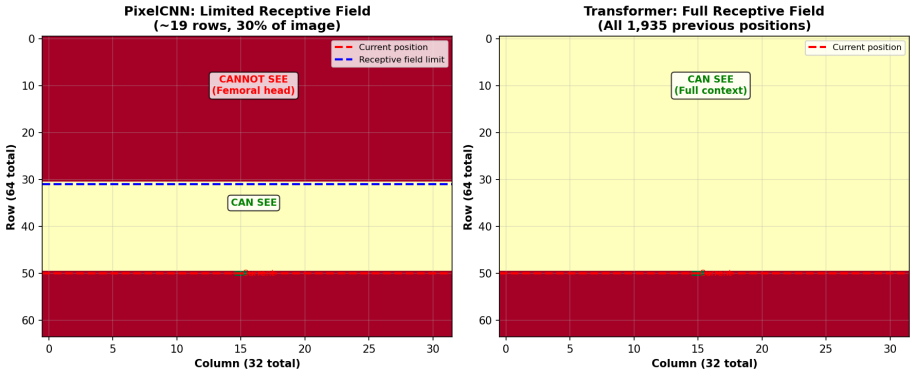


Figure 17: Receptive field comparison

A.2 Receptive Field Visualization

Figure A.2: Visual comparison of PixelCNN's limited receptive field (~19 rows) versus Transformer's full attention over all 2,048 code positions.

A.3 Transformer Sample Grid

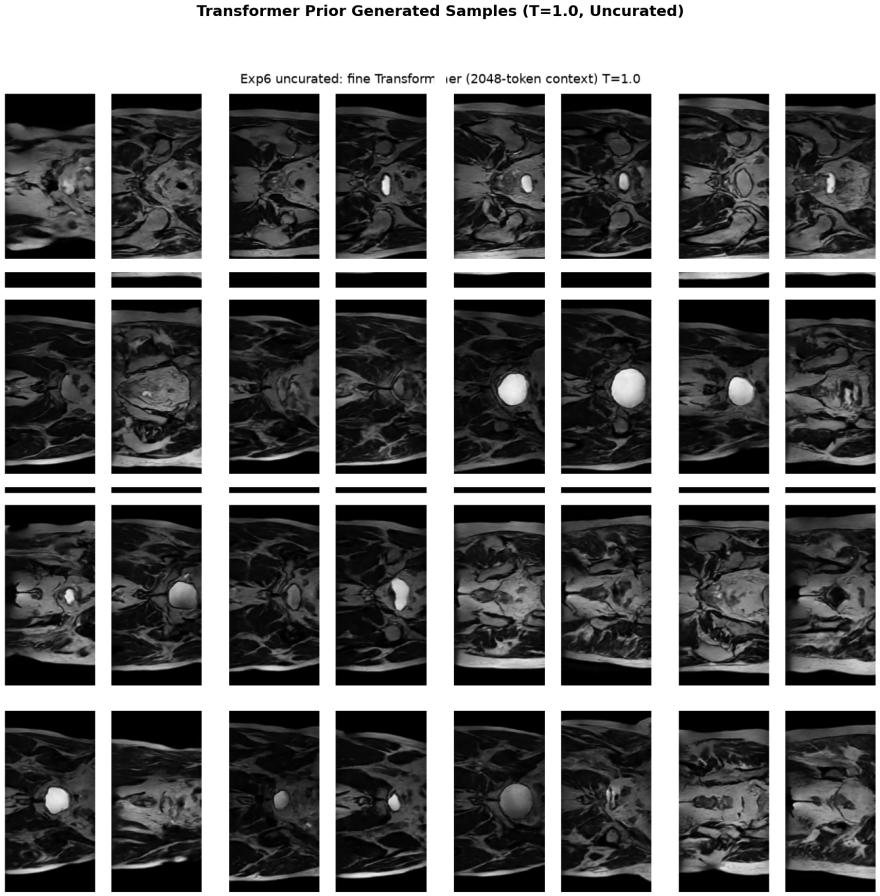


Figure 18: Transformer sample grid

Figure A.3: Extended grid of Transformer-generated samples (32 samples) demonstrating diversity and anatomical consistency across multiple generations.

A.4 Distribution Analysis (Experiment 6)

Figure A.4: Detailed statistical analysis from Experiment 6 (Transformer prior) showing intensity distributions, bright pixel fractions, and pairwise distance metrics.

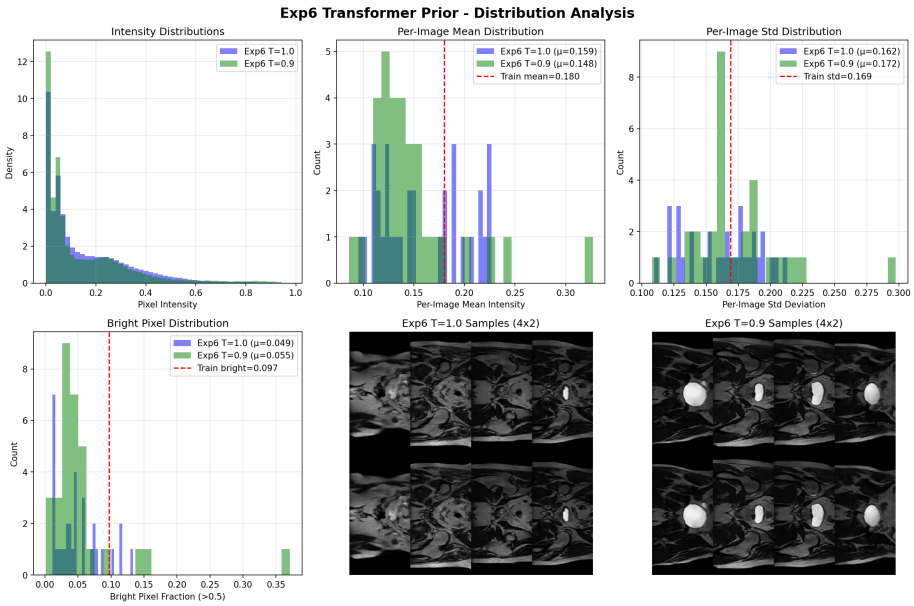


Figure 19: Distribution analysis

Appendix B: Complete Feasibility Review (As Submitted)

This is the exact 1-page feasibility review submitted before experiments began.

COMP3710 Project 2.4: Feasibility Review

The Fidelity–Memorisation–Privacy Dilemma in Synthetic Medical Imaging
Hard VQ-VAE with PixelCNN Prior | DCGAN Baseline

---

1. User Need, Scope & Acceptance Criteria

Synthetic medical images could support research when real patient data is limited or cannot be shared due to privacy constraints. However, generated images must preserve anatomical structure without simply reproducing training examples. This project investigates these trade-offs using 2D HipMRI slices, comparing a VQ-VAE with a PixelCNN prior against a DCGAN baseline.

The project will assess the following criteria:

- **Reconstruction fidelity:** achieve mean SSIM 0.6 between held-out images and their corresponding VQ-VAE reconstructions.
- **Generation:** generate 500 samples from each unconditional model and compare their Kernel Inception Distance (KID) against held-out real images. Lower KID indicates greater similarity in feature distributions.
- **Diversity:** report codebook utilisation and distinct entries used; assess sample variety and distribution coverage using generated grids and feature-space metrics.
- **Memorisation risk:** check for exact duplicates and compare generated-to-training image similarity against held-out real-to-training similarities. Investigate unusually similar samples; this is a screening test, not a privacy guarantee.
- **Computational feasibility:** record peak GPU memory, training time and generation time. Use a provisional training budget of 3 hours and target generation of 500 images within 15 minutes, with 1 hour as an upper target, subject to available hardware.

2. Model Choice & Course Concepts

The main model is a VQ-VAE trained on the provided HipMRI dataset. Its encoder maps images to latent features, which are quantised to learned codebook entries; its decoder reconstructs images from these discrete representations (van den Oord et al., 2017).

A VQ-VAE alone does not learn how to sample new latent grids. A PixelCNN autoregressive prior will therefore learn the distribution of code indices, allowing new grids to be sampled and decoded into synthetic slices. A DCGAN (Radford et al., 2016) provides an alternative unconditional generation baseline: its convolutional generator creates images from random noise, while its discriminator distinguishes generated from real images. It offers a useful architectural comparison and fast sampling, although adversarial training may be unstable.

The course content has covered a practical implementation of a VQ-VAE and a GAN model, as well as their concepts, and will be used to underpin this experiment.

3. Preliminary Feasibility Evidence

Baseline configuration. The VQ-VAE follows the architecture and settings of van den Oord et al. (2017): two 4×4 stride-2 convolutions (a 64×32 latent grid), two residual blocks, a 512-entry codebook of dimension 64, =0.25, straight-through gradients and Adam at 2×10 , with the width reduced to 128 channels. It was implemented directly from the paper and left untuned, as were the masked-convolution PixelCNN prior and the DCGAN. A previous-year COMP3710 VQ-VAE on this dataset (PatternAnalysis-2025, s4858817) is

cited only as external evidence: it reconstructs without a generative prior and reached validation SSIM 0.757, with codebook perplexity rising from ~5 to ~100 over ~50 epochs.

Data audit (Rangpur, keras_slices_data). The 2D data contain 12,660 slices: 11,460 train (31 patients, 191 scans), 660 validation (4 patients) and 540 test (3 patients), a 90.5/5.2/4.3% slice split with zero patient overlap between splits. Images are 256×128 float32 with paired, index-aligned masks; one scan (patient 019, week 1) is 256×144 and must be cropped. Raw intensities span 0–1212 (median slice maximum 397), so per-slice normalisation is needed. Masks carry six labels; the prostate (label 5) is only ~0.4% of pixels. Loading runs at ~230 slices/s; the full set is ~1.7 GB in RAM.

GPU smoke test (A100-40GB, 5 epochs). The VQ-VAE (1.18M parameters) trained at 6.9 s/epoch and 3.39 GB peak, reaching held-out mean SSIM 0.624 (78% of slices 0.6; MSE 0.0043), though reconstructions are blurry. Only 13/512 codes were used (perplexity 6.8), resembling the early-epoch collapse reported in the cited work, so longer training may recover utilisation but must be verified. A 2.13M-parameter PixelCNN on the frozen codes cut cross-entropy from 6.24 (uniform) to 0.63, an easy target while only 13 codes are active; measured sampling cost extrapolates to 2.8 min (fp32; 1.4 min bf16) for 500 images, inside the 15-minute target. A 3-epoch DCGAN (3.2M/2.8M parameters, 3.6 s/epoch, 1.46 GB) generates 500 images in 0.12 s but gives banded, near-identical samples (pixel std 0.067 vs 0.130 real). The pipeline fits the budget; sample quality is the open risk.

Feasibility figure

Figure 20: Feasibility figure

Fig. 1. Real vs VQ-VAE reconstruction pairs (blurry, SSIM 0.62); DCGAN samples (banded, near-identical).

4. Experimental Plan & Next Steps

VQ-VAE: - Retrain for 50–100 epochs with validation-SSIM checkpointing and per-epoch perplexity, testing whether the 13/512 collapse is transient. - If utilisation stays low: EMA codebook updates, k-means initialisation, dead-code reset; sweep K {256,512}, D {64,128} and . - To reduce blur: add an L1/SSIM loss and compare 4× (2048 tokens) with 8× downsampling (512 tokens: ~4× cheaper PixelCNN sampling, possible detail loss). Select on validation SSIM, MSE, utilisation.

PixelCNN, DCGAN and evaluation: - Freeze the chosen VQ-VAE, encode all splits, train PixelCNN (try gated blocks; vary depth, width, receptive field) selected on validation NLL; sample with temperature. - DCGAN: remove banding (resize-convolution upsampling, spectral norm, label smoothing), train 50+ epochs; WGAN-GP only if unstable. - 500 samples per model: KID vs held-out

real images, codebook use, fixed-seed grids, nearest-training-neighbour audit. Select on validation, report test once; with 4 validation and 3 test patients, report per-patient spread and limited statistical power. Log time and peak GPU memory.

5. Risks, Budget & Fallback

Risks: Blurry reconstructions, codebook under-utilisation, poor generated samples, DCGAN instability, memorisation and limited computing resources.

- **Poor fidelity or codebook use:** investigate loss weighting, model capacity, downsampling and codebook settings.
- **Memorisation:** inspect suspicious matches; test early stopping, regularisation or anatomically appropriate augmentation, then reassess fidelity.
- **Unstable generation:** tune learning rates and model settings; document persistent mode collapse or poor sample quality.
- **Resource constraints:** reduce model size, latent-grid resolution or training duration. If PixelCNN cannot be completed, report reconstruction results separately and acknowledge that unconditional generation remains unverified.

References

van den Oord, A., Vinyals, O. & Kavukcuoglu, K. (2017). *Neural Discrete Representation Learning*. NeurIPS.

Radford, A., Metz, L. & Chintala, S. (2016). *Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks*. ICLR.